

Fair Locks and Semaphores in Kotlin Coroutines

Никита Коваль, JetBrains
Роман Елизаров, JetBrains

ITMO 2020

Locks and Semaphores

Mutex = **M**utual **E**xclusion, at most 1 thread is in the critical section

Locks and Semaphores

Mutex = **M**utual **E**xclusion, at most 1 thread is in the critical section

```
val m = Mutex()
```

```
m.lock()
```

```
[critical section]
```

```
m.unlock()
```

```
m.lock()
```

```
[critical section]
```

```
m.unlock()
```

Locks and Semaphores

Mutex = **M**utual **E**xclusion, at most 1 thread is in the critical section

```
val m = Mutex()
```

```
m.lock()  
[critical section]  
m.unlock()
```

transfer
the permit

```
m.lock()  
[critical section]  
m.unlock()
```

Locks and Semaphores

Semaphore = at most **K** threads are in the critical section

Mutex = Semaphore(**permits = 1**)

Locks and Semaphores

Semaphore = at most **K** threads are in the critical section

Mutex = Semaphore(**permits** = 1)

Semaphore algorithm is
the start of this project!

Blocking Calls via Future

Threads, coroutines, continuations... Need a simple abstraction!

Blocking Calls via Future

Threads, coroutines, continuations... Need a simple abstraction!

```
interface Future<T> {  
    fun await(): T  
}
```


Blocking Calls via Future

Threads, coroutines, continuations... Need a simple abstraction!

```
interface Future<T> {  
    fun await(): T  
}
```

```
class FutureImmediate<T>(  
    val res: T  
) {  
    fun await(): T = res  
}
```

Blocking Calls via Future


Threads, coroutines, continuations... Need a simple abstraction!

```
interface Future<T> {  
    fun await(): T  
}
```


```
class FutureImmediate<T>(  
    val res: T  
) {  
    fun await(): T = res  
}
```

```
class FutureSuspended<T> {  
    var state: T? = null  
  
    fun await(): T {  
        while (state == null) {}  
        return state  
    }  
  
    fun complete(value: T) {  
        state = value  
    }  
}
```

State Machine for FutureSuspended



State Machine for FutureSuspended



Further **FutureSuspended** updates
will be shown via such diagrams

Semaphore API

```
class Semaphore(permits: Int) {  
    fun acquire(): Unit  
    fun release()  
}
```

Semaphore API

```
class Semaphore(permits: Int) {  
    fun acquire(): Future<Unit>  
    fun release()  
}
```

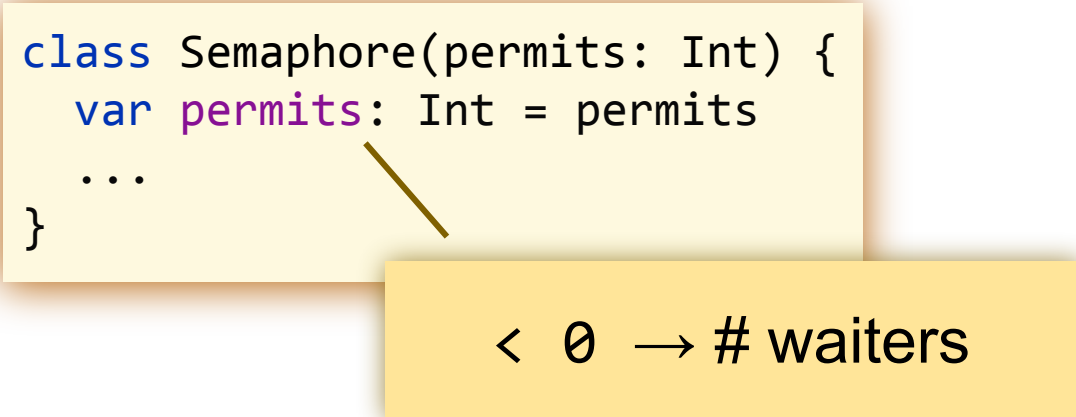
Blocking by design

Straightforward Semaphore Implementation

```
class Semaphore(permits: Int) {  
    var permits: Int = permits  
    ...  
}
```

Straightforward Semaphore Implementation

```
class Semaphore(permits: Int) {  
    var permits: Int = permits  
    ...  
}
```



$< 0 \rightarrow \# \text{ waiters}$

Straightforward Semaphore Implementation

```
class Semaphore(permits: Int) {  
    var permits: Int = permits  
    ...  
}
```

```
fun acquire(): Future<Unit> {  
    val p = FAA(&permits, -1)  
    if p > 0:  
        return FutureImmediate(Unit)  
    else:  
        return suspend()  
}
```

Creates a new FutureSuspended and puts it into the waiting queue

Straightforward Semaphore Implementation


```
class Semaphore(permits: Int) {  
    var permits: Int = permits  
    ...  
}
```

Retrieves the first waiter
and completes it

```
fun acquire(): Future<Unit> {  
    val p = FAA(&permits, -1)  
    if p > 0:  
        return FutureImmediate(Unit)  
    else:  
        return suspend()  
}
```


```
fun release() {  
    val p = FAA(&permits, +1)  
    if p >= 0: return  
    resume(Unit)  
}
```

SegmentQueueSynchronizer




An infinite array with
indices for the next
addition and retrieval

SegmentQueueSynchronizer




```
fun suspend(): Future<T> {  
    val f = FutureSuspended<T>()  
    val i = FAA(&enqIdx, +1)  
    // store f into w[i]  
}
```

SegmentQueueSynchronizer




```
fun suspend(): Future<T> {  
    val f = FutureSuspended<T>()  
    val i = FAA(&enqIdx, +1)  
    // store f into w[i]  
}
```

SegmentQueueSynchronizer




```
fun suspend(): Future<T> {  
    val f = FutureSuspended<T>()  
    val i = FAA(&enqIdx, +1)    i:0  
    // store f into w[i]  
}
```

SegmentQueueSynchronizer



```
fun suspend(): Future<T> {  
    val f = FutureSuspended<T>()  
    val i = FAA(&enqIdx, +1)  i:0  
    // store f into w[i]  
}
```


SegmentQueueSynchronizer



```
fun suspend(): Future<T> {  
    val f = FutureSuspended<T>()  
    val i = FAA(&enqIdx, +1)  
    // store f into w[i]  
}
```

```
fun resume(value: T) {  
    val i = FAA(&deqIdx, +1)  
    // complete the future  
    // located in w[i]  
}
```



SegmentQueueSynchronizer



```
fun suspend(): Future<T> {  
    val f = FutureSuspended<T>()  
    val i = FAA(&enqIdx, +1)  
    // store f into w[i]  
}
```

```
fun resume(value: T) {  
    val i = FAA(&deqIdx, +1) i:0  
    // complete the future  
    // located in w[i]  
}
```

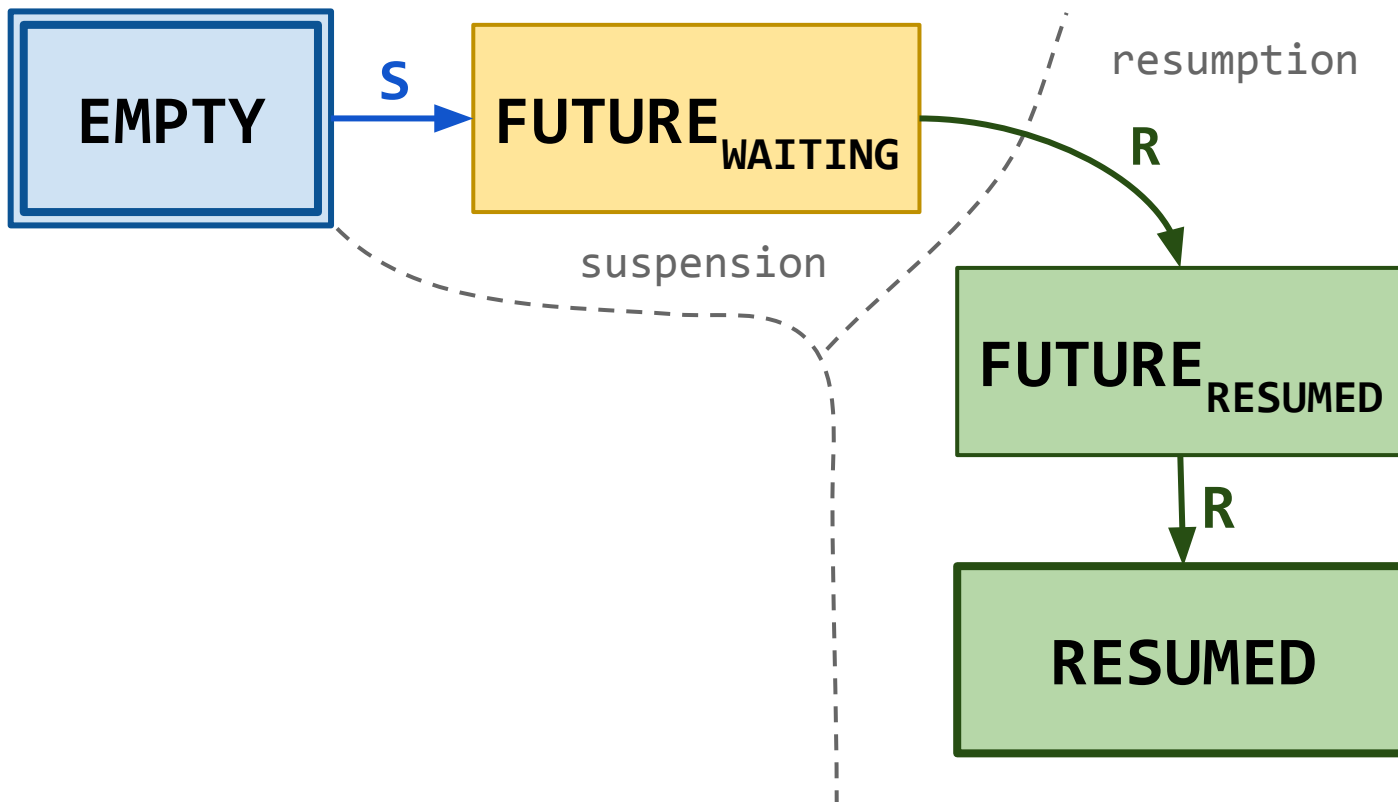
SegmentQueueSynchronizer



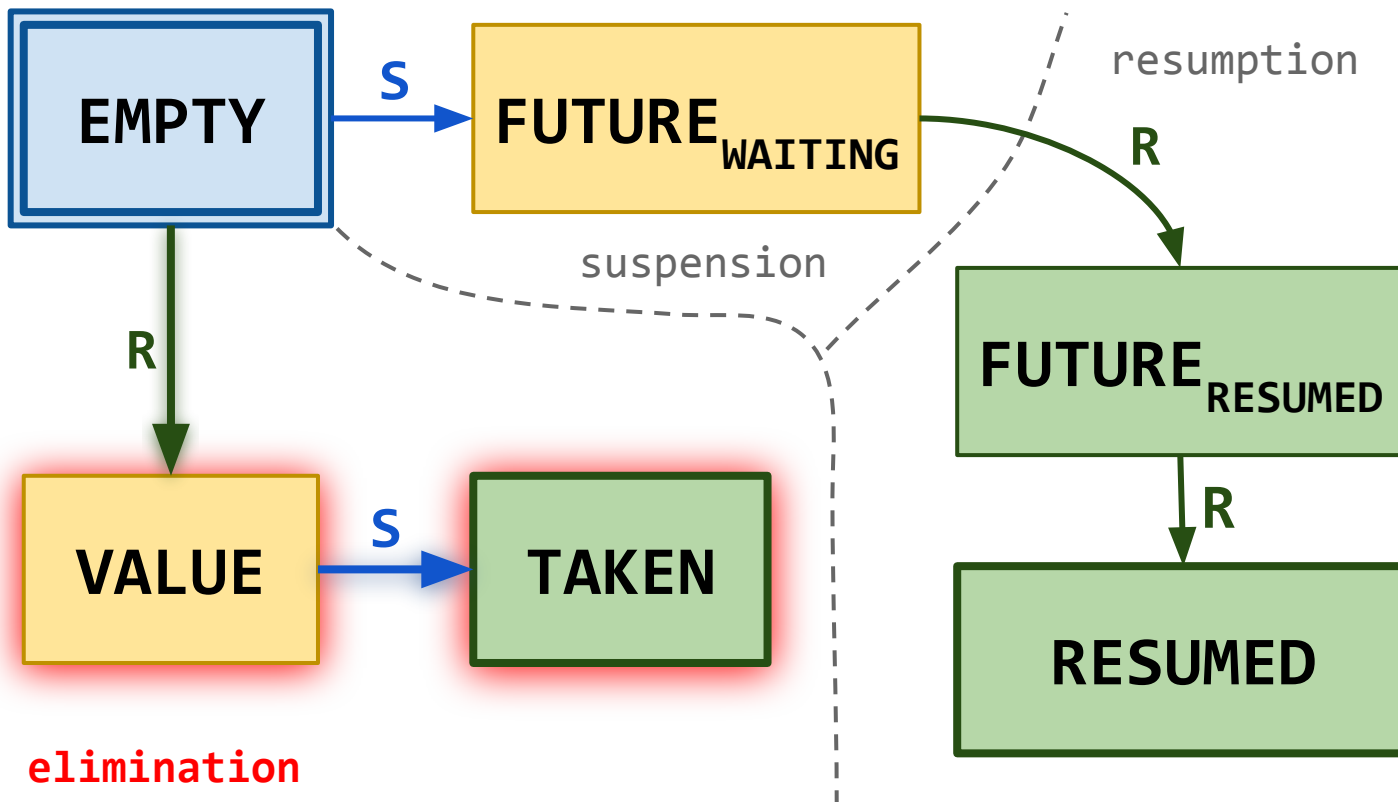
```
fun suspend(): Future<T> {  
    val f = FutureSuspended<T>()  
    val i = FAA(&enqIdx, +1)  
    // store f into w[i]  
}
```

```
fun resume(value: T) {  
    val i = FAA(&deqIdx, +1) i:0  
    // complete the future  
    // located in w[i]  
}
```


Cell Life-Cycle




Cell Life-Cycle



Infinite Array Implementation




Infinite Array Implementation



```
fun suspend(): Future<T> {  
    val t = tail  
    val i = FAA(&enqIdx, +1)  
    val s = findSegment(start = t,  
                        id = i / M)  
    moveTailForward(t)  
    // w[i] is s[i % M]  
}
```

Infinite Array Implementation



```
fun suspend(): Future<T> {  
    val t = tail  
    val i = FAA(&enqIdx, +1)  
    val s = findSegment(start = t,  
                        id = i / M)  
    moveTailForward(t)  
    // w[i] is s[i % M]  
}
```

```
fun resume(value: T) {  
    val t = head  
    val i = FAA(&deqIdx, +1)  
    val s = findSegment(start = t,  
                        id = i / M)  
    moveHeadForward(t)  
    // w[i] is s[i % M]  
}
```

Extend Semaphore with tryAcquire

```
fun acquire(): Future<Unit> {  
    val p = FAA(&permits, -1)  
    if p > 0:  
        return FutureImmediate(Unit)  
    else:  
        return suspend()  
}
```

```
fun release() {  
    val p = FAA(&permits, +1)  
    if p >= 0: return  
    resume(Unit)  
}
```


Extend Semaphore with tryAcquire

```
fun acquire(): Future<Unit> {  
    val p = FAA(&permits, -1)  
    if p > 0:  
        return FutureImmediate(Unit)  
    else:  
        return suspend()  
}
```

```
fun release() {  
    val p = FAA(&permits, +1)  
    if p >= 0: return  
    resume(Unit)  
}
```


```
fun tryAcquire(): Boolean = while(true) {  
    val p = permits  
    if p <= 0: return false  
    if CAS(&permits, p, p - 1): return true  
}
```

Is this tryAcquire correct?

```
fun acquire(): Future<Unit> {  
    val p = FAA(&permits, -1)  
    if p > 0:  
        return FutureImmediate(Unit)  
    else:  
        return suspend()  
}
```

```
fun tryAcquire(): Boolean {  
    while(true) {  
        val p = permits  
        if p <= 0: return false  
        if CAS(&permits, p, p - 1):  
            return false  
    }  
}
```

```
fun release() {  
    val p = FAA(&permits, +1)  
    if p >= 0: return  
    resume(Unit)  
}
```



```
val s = Semaphore(1); s.acquire()
```

```
s.release()  
s.tryAcquire()  
s.acquire()
```

```
s.acquire():  
    1. dec permits  
    2. suspend
```


Is this tryAcquire correct?

```
fun acquire(): Future<Unit> {  
    val p = FAA(&permits, -1)  
    if p > 0:  
        return FutureImmediate(Unit)  
    else:  
        return suspend()  
}
```

```
fun tryAcquire(): Boolean {  
    while(true) {  
        val p = permits  
        if p <= 0: return false  
        if CAS(&permits, p, p - 1):  
            return false  
    }  
}
```

```
fun release() {  
    val p = FAA(&permits, +1)  
    if p >= 0: return  
    resume(Unit)  
}
```

enqIdx



Semaphore

permits = -1

```
val s = Semaphore(1); s.acquire()
```

```
s.release()  
s.tryAcquire()  
s.acquire()
```


```
s.acquire():  
    1. dec permits  
    2. suspend
```

Is this tryAcquire correct?

```
fun acquire(): Future<Unit> {  
    val p = FAA(&permits, -1)  
    if p > 0:  
        return FutureImmediate(Unit)  
    else:  
        return suspend()  
}
```

```
fun tryAcquire(): Boolean {  
    while(true) {  
        val p = permits  
        if p <= 0: return false  
        if CAS(&permits, p, p - 1):  
            return false  
    }  
}
```

```
fun release() {  
    val p = FAA(&permits, +1)  
    if p >= 0: return  
    resume(Unit)  
}
```



Semaphore

permits = 0

```
val s = Semaphore(1); s.acquire()
```

```
s.release() :done  
s.tryAcquire()  
s.acquire()
```


```
s.acquire():  
    1. dec permits  
    2. suspend
```

Is this tryAcquire correct?

```
fun acquire(): Future<Unit> {  
    val p = FAA(&permits, -1)  
    if p > 0:  
        return FutureImmediate(Unit)  
    else:  
        return suspend()  
}
```

```
fun tryAcquire(): Boolean {  
    while(true) {  
        val p = permits  
        if p <= 0: return false  
        if CAS(&permits, p, p - 1):  
            return false  
    }  
}
```

```
fun release() {  
    val p = FAA(&permits, +1)  
    if p >= 0: return  
    resume(Unit)  
}
```



```
val s = Semaphore(1); s.acquire()
```

```
s.release() :done  
s.tryAcquire():f  
s.acquire()
```


```
s.acquire():  
    1. dec permits  
    2. suspend
```

Is this tryAcquire correct?

```
fun acquire(): Future<Unit> {  
    val p = FAA(&permits, -1)  
    if p > 0:  
        return FutureImmediate(Unit)  
    else:  
        return suspend()  
}
```

```
fun tryAcquire(): Boolean {  
    while(true) {  
        val p = permits  
        if p <= 0: return false  
        if CAS(&permits, p, p - 1):  
            return false  
    }  
}
```

```
fun release() {  
    val p = FAA(&permits, +1)  
    if p >= 0: return  
    resume(Unit)  
}
```



Semaphore

permits = -1


```
val s = Semaphore(1); s.acquire()
```

```
s.release() :done  
s.tryAcquire():f  
s.acquire() :done
```

```
s.acquire():  
    1. dec permits  
    2. suspend
```

Is this tryAcquire correct?

release was intended to give a permit to a concurrent **acquire**, but gave it to **acquire** that *happens after* it

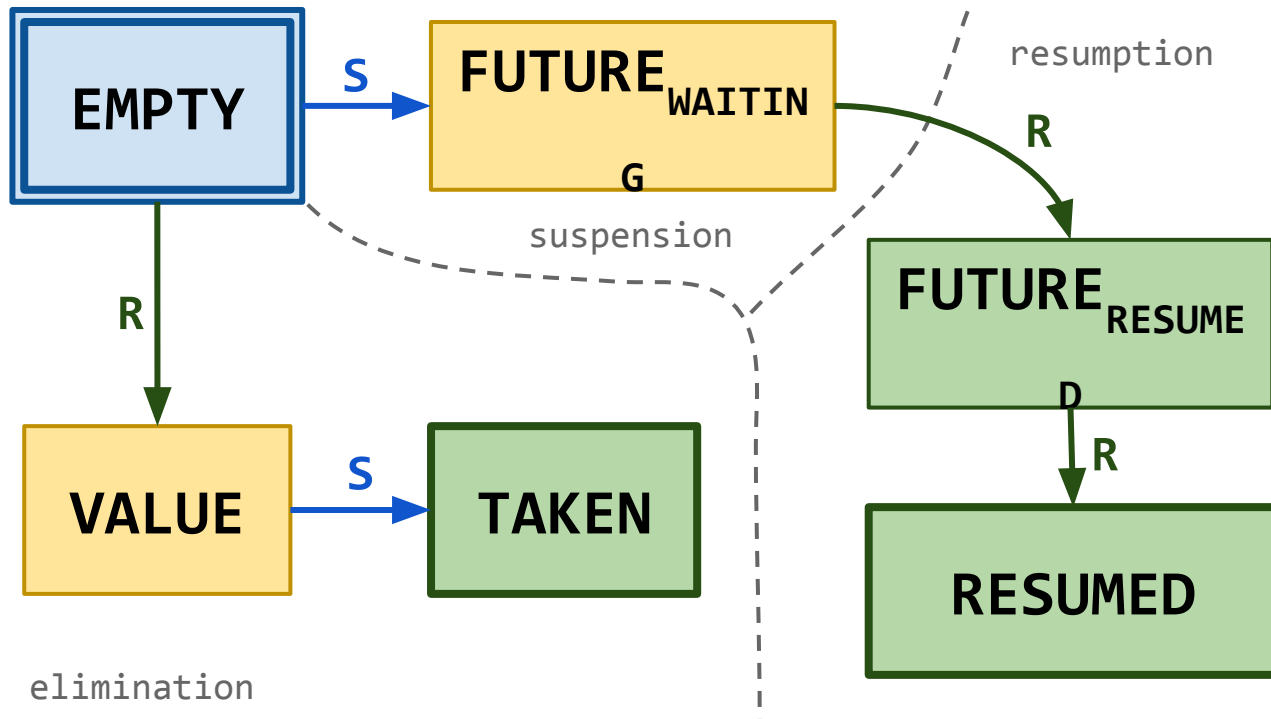


```
val s = Semaphore(1); s.acquire()
```

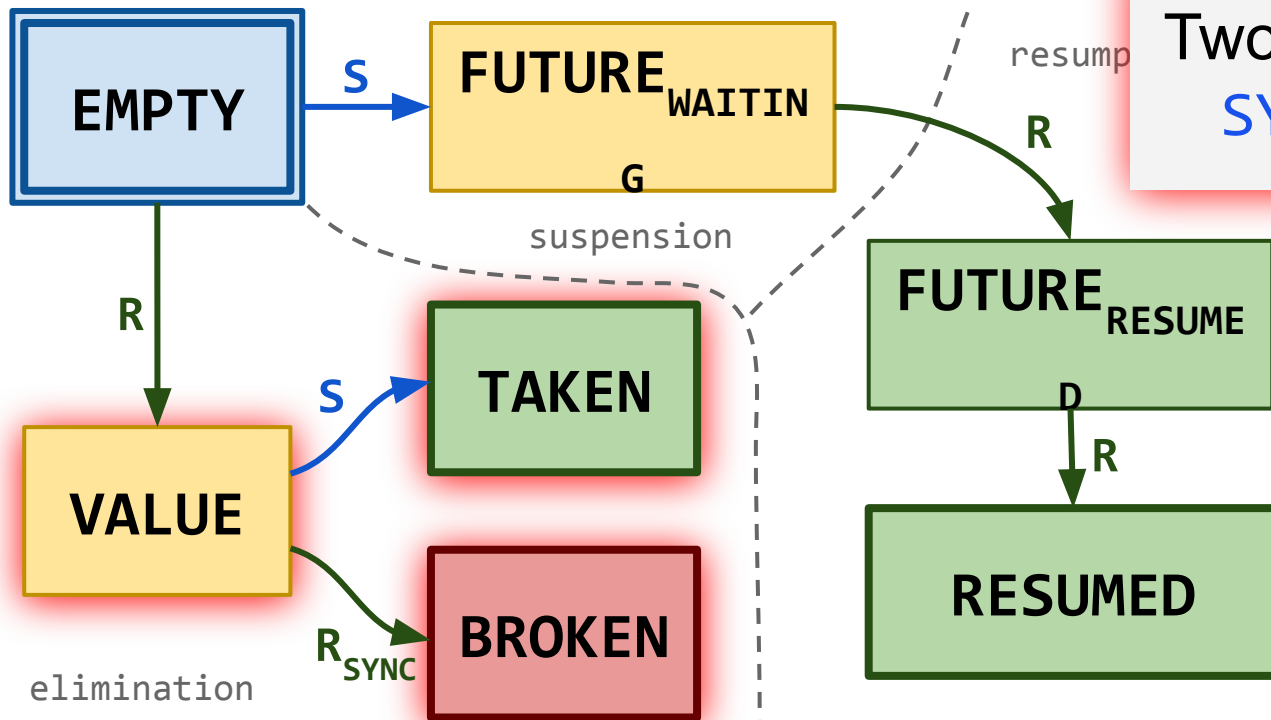
```
s.release() :done  
s.tryAcquire():f  
s.acquire() :done
```

```
s.acquire():  
  1. dec permits  
  2. suspend
```

Extend Semaphore with tryAcquire



Extend Semaphore with tryAcquire



Two resume modes:
SYNC and **ASYNC**

Extend Semaphore with tryAcquire

```
class SegmentQueueSynchronizer<T> {  
    // can fail on elimination  
    fun resume(): Boolean { ... }  
  
    // returns null on broken cell  
    fun suspend(): Future<T>? { ... }  
}
```

Extend Semaphore with tryAcquire

```
class SegmentQueueSynchronizer<T> {  
    // can fail on elimination  
    fun resume(): Boolean { ... }  
  
    // returns null on broken cell  
    fun suspend(): Future<T>? { ... }  
}
```

```
fun acquire(): Future<Unit> {  
    while(true) {  
        val p = FAA(&permits, -1)  
        if p > 0:  
            return FutureImmediate(Unit)  
        val f = suspend()  
        if f != null: return f  
    }  
}
```

```
fun release() {  
    while(true) {  
        val p = FAA(&permits, +1)  
        if p >= 0: return  
        val done = resume(Unit)  
        if done: return  
    }  
}
```


Cancellable FutureSuspended

- Cancellation (abortability) is natural for blocking primitives
- Moreover, it is a built-in feature in Kotlin Coroutines

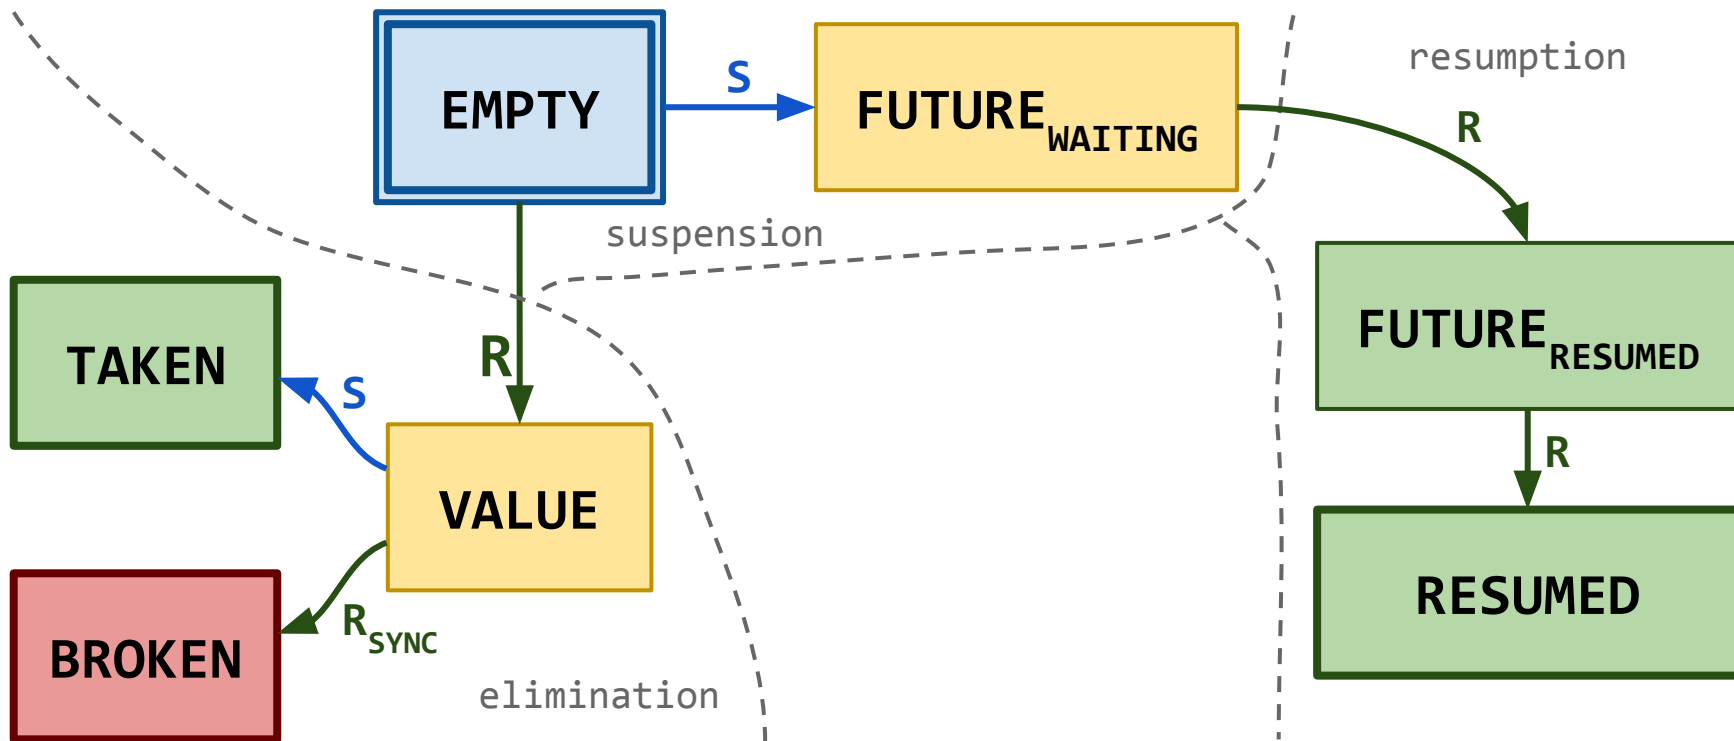
Cancellable FutureSuspended

- Cancellation (abortability) is natural for blocking primitives
- Moreover, it is a built-in feature in Kotlin Coroutines

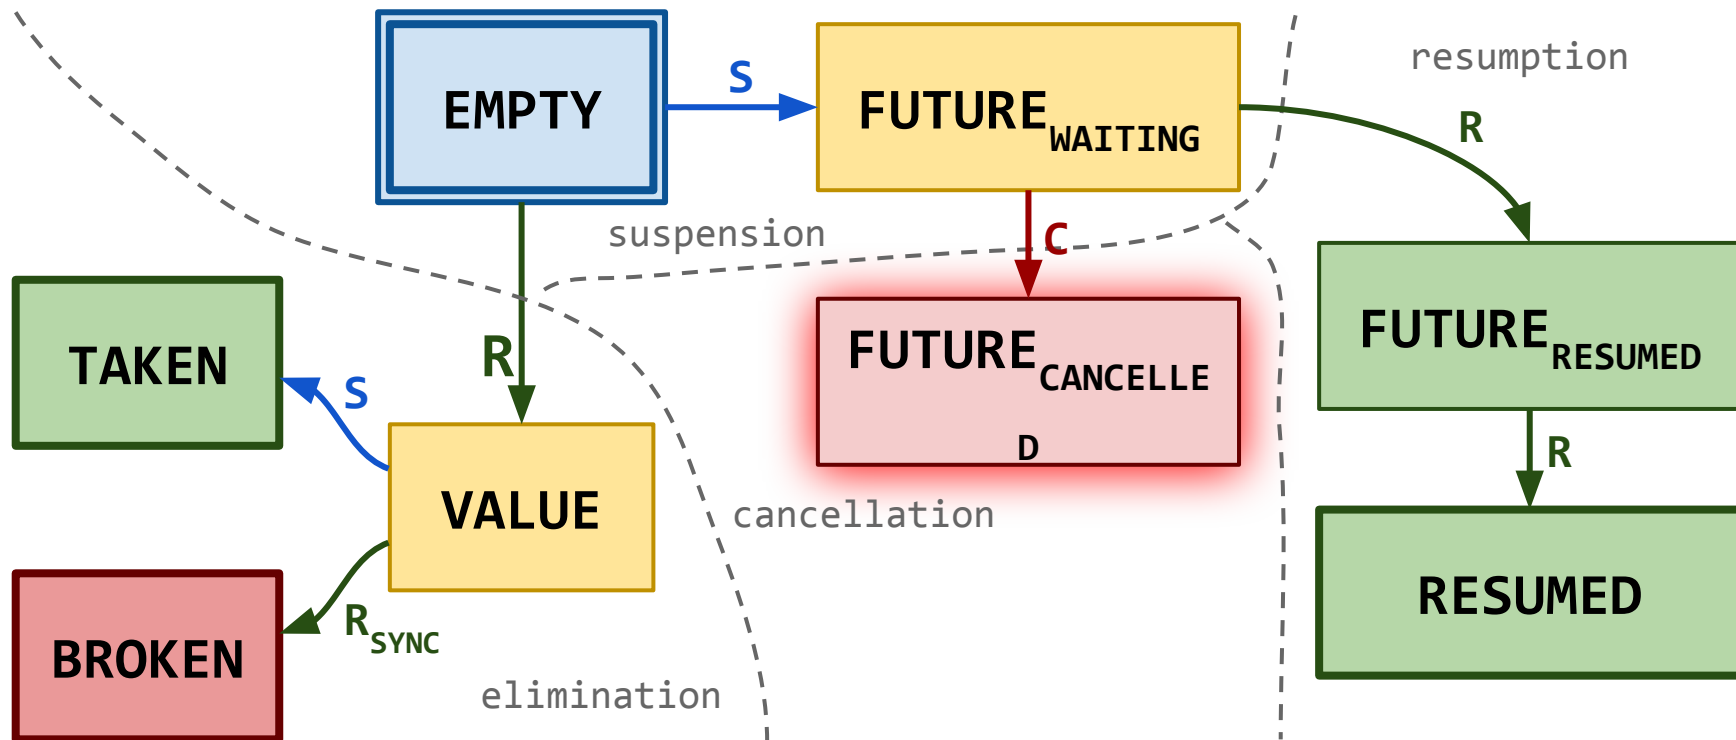
```
class FutureSuspended<T> {  
    // returns  $\perp$  if cancelled  
    fun await(): T? { ... }  
    fun complete(value: T) { ... }  
    fun cancel(): Boolean { ... }  
  
    fun handleCancellation() {  
        // TODO: Implement me, please!  
    }  
}
```



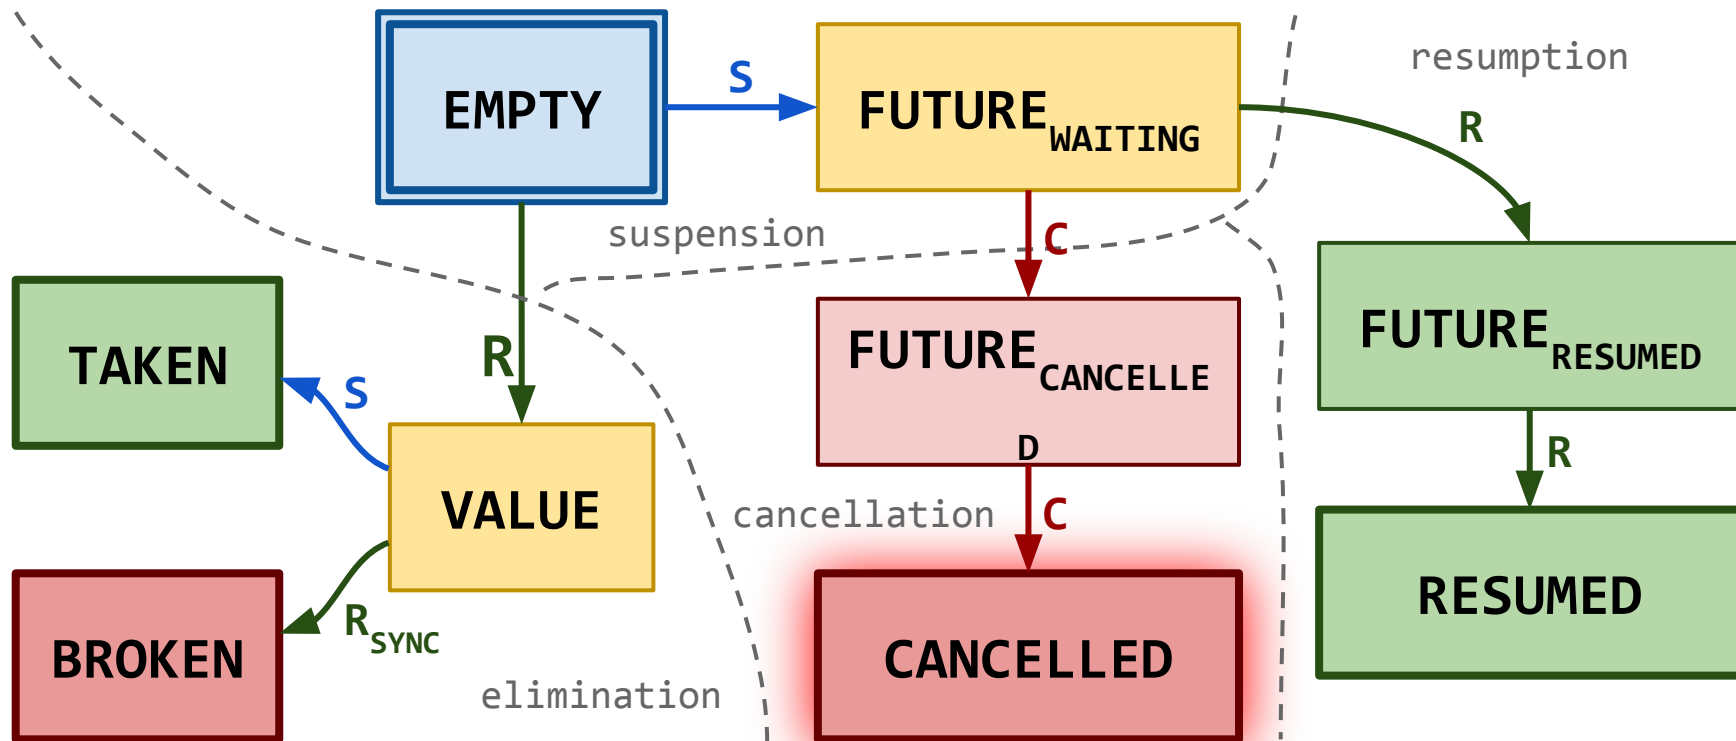
Cancellation Support in the Cell Life-Cycle



Cancellation Support in the Cell Life-Cycle



Cancellation Support in the Cell Life-Cycle



Cancellation Support in the Cell Life-Cycle

```
class FutureSuspended<T> {  
    ...  
  
    fun handleCancellation() {  
        // move the cell to  
        // `CANCELLED` state  
    }  
}
```


BROKEN

elimination

FUTURE_{WAITING}

FUTURE_{CANCELLE}

CANCELLED



Cancellation Support in the Cell Life-Cycle

```
class FutureSuspended<T> {  
    ...  
  
    fun handleCancellation()  
        // move the cell to  
        // `CANCELLED` state  
    }  
}
```

FUTURE

resumption

```
class SegmentQueueSynchronizer<T> {  
    ...  
    // fails if the next  
    // waiter is cancelled  
    fun resume(): Boolean {  
        ...  
    }  
}
```

BROKEN

R_{SYNC}


elimination

Cancellation Support in General

- `handleCancellation` moves the cell state to **CANCELLED** to avoid memory leaks
- Can we remove the cells themselves for the same reason?

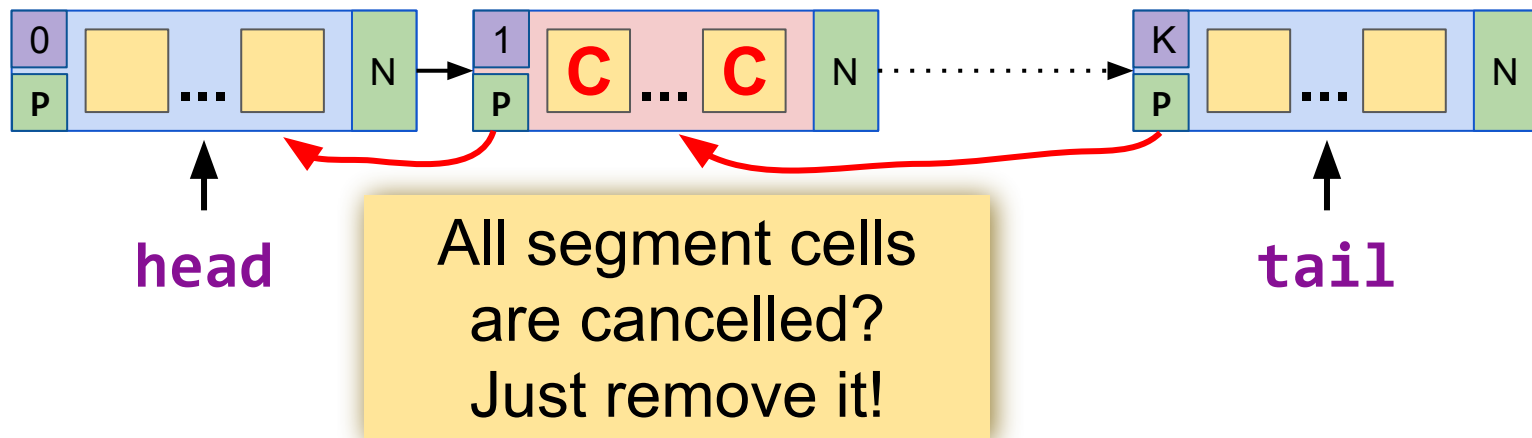
Cancellation Support and Memory Leaks

- `handleCancellation` moves the cell state to **CANCELLED** to avoid memory leaks
- Can we remove the cells themselves for the same reason?



Cancellation Support and Memory Leaks

- `handleCancellation` moves the cell state to **CANCELLED** to avoid memory leaks
- Can we remove the cells themselves for the same reason?



Semaphore with Cancellation

Our implementation is already correct!

```
fun acquire(): Future<Unit> {  
    while(true) {  
        val p = FAA(&permits, -1)  
        if p > 0:  
            return FutureImmediate(Unit)  
        val f = suspend()  
        if f != null: return f  
    }  
}
```

```
fun release() {  
    while(true) {  
        val p = FAA(&permits, +1)  
        if p >= 0: return  
        // can fail due to cancellation,  
        // adjusted `permits` then  
        val done = resume(Unit)  
        if done: return  
    }  
}
```

Evaluation

- Google Cloud machine with 96 CPUs
- Comparison against standard Java implementations:
 - `j.u.c.Semaphore` and `j.u.c.ReentrantLock`


Evaluation

- Google Cloud machine with 96 CPUs
- Comparison against standard Java implementations:
 - `j.u.c.Semaphore` and `j.u.c.ReentrantLock`

```
fun operation()  
{  
    semaphore.acquire()  
  
    doSomeGeomDistrWork()  
  
    semaphore.release()  
}
```





Let's run this code
on different numbers
of threads!

Evaluation



- Java ReentrantLock
- Java Semaphore
- SYNC + SIMPLE
- ASYNC + SIMPLE
- SYNC + SMART
- ASYNC + SMART

Evaluation



Let's use `SegmentQueueSynchronizer`
for primitives other than `Semaphore`!

CountDownLatch

Allows to wait until several operations are completed

```
class CountDownLatch(count: Int) {  
    fun countDown() { ... }  
    fun await(): Future<Unit> { ... }  
}
```

CountDownLatch

resume_mode = **ASYNC**

```
class CountDownLatch(count: Int) {  
    val count = count  
    val waiters = 0  
    ...  
}
```

CountDownLatch

resume_mode = **ASYNC**

```
class CountDownLatch(count: Int) {  
    val count = count  
    val waiters = 0  
    ...  
}
```

```
fun countdown() {  
    val c = FAA(&count, -1)  
    if c <= 0: resumeWaiters()  
}
```

CountDownLatch

resume_mode = ASYNC

```
class CountDownLatch(count: Int) {  
    val count = count  
    val waiters = 0  
    ...  
}
```

```
fun countDown() {  
    val c = FAA(&count, -1)  
    if c <= 0: resumeWaiters()  
}  
  
fun resumeWaiters() = while (true) {  
    val w = waiters  
    if w & DONE_BIT != 0: return  
    if CAS(&waiters, w, w & DONE_BIT) {  
        repeat(w) { resume(Unit) }  
    }  
}
```

CountDownLatch

resume_mode = ASYNC

```
class CountDownLatch(count: Int) {  
    val count = count  
    val waiters = 0  
    ...  
}
```

```
fun await(): Future<Unit> {  
    if count <= 0:  
        return FutureImmediate(Unit)  
    val w = FAA(&waiters, +1)  
    if w & DONE_BIT != 0:  
        return FutureImmediate(Unit)  
    return suspend()  
}
```

```
fun countDown() {  
    val c = FAA(&count, -1)  
    if c <= 0: resumeWaiters()  
}
```

```
fun resumeWaiters() = while (true) {  
    val w = waiters  
    if w & DONE_BIT != 0: return  
    if CAS(&waiters, w, w & DONE_BIT) {  
        repeat(w) { resume(Unit) }  
    }  
}
```


Pools

resume_mode = ASYNC

```
class BlockingPool() {  
  // < 0 => # waiters  
  var elements = 0  
  ...  
}
```

Pools

resume_mode = **ASYNC**

```
class BlockingPool() {  
    // < 0 => # waiters  
    var elements = 0  
    ...  
}
```

```
fun put(element: T) {  
    val e = FAA(&elements, 1)  
    if e < 0 {  
        resume(element)  
    } else {  
        insertIntoPool(element)  
    }  
}
```

Pools

resume_mode = **ASYNC**

```
class BlockingPool() {  
    // < 0 => # waiters  
    var elements = 0  
    ...  
}
```

```
fun put(element: T) {  
    val e = FAA(&elements, 1)  
    if e < 0 {  
        resume(element)  
    } else {  
        insertIntoPool(element)  
    }  
}
```

```
fun retrieve(): Future<T> {  
    val e = FAA(&elements, -1)  
    if (e > 0) {  
        val elem = retrieveFromPool()  
        return FutureImmediate(elem)  
    } else {  
        return suspend()  
    }  
}
```

Pools

```
resume_mode = ASYNC
```

```
class BlockingPool() {  
    // < 0 => # waiters  
    var elements = 0  
    ...  
}
```

```
fun put(element: T) {  
    val e = FAA(&elements, 1)  
    if e < 0 {  
        resume(element)  
    } else {  
        insertIntoPool(element)  
    }  
}
```

```
fun retrieve(): Future<T> {  
    val e = FAA(&elements, -1)  
    if (e > 0) {  
        val elem = retrieveFromPool()  
        return FutureImmediate(elem)  
    } else {  
        return suspend()  
    }  
}
```

insertIntoPool and **retrieveFromPool**
can use any data structure under the hood,
e.g. queue, stack, or just bag.

Pools

resume_mode = **ASYNC**

```
class BlockingPool() {  
    // < 0 => # waiters  
    var elements = 0  
    ...  
}
```

Smart cancellation mode can
be used here as well

```
fun put(element: T) {  
    val e = FAA(&elements, 1)  
    if e < 0 {  
        resume(element)  
    } else {  
        insertIntoPool(element)  
    }  
}
```

```
...(): Future<T> {  
    val e = FAA(&elements, -1)  
    if (e > 0) {  
        val elem = retrieveFromPool()  
        return FutureImmediate(elem)  
    } else {  
        return suspend()  
    }  
}
```

insertIntoPool and **retrieveFromPool**
can use any data structure under the hood,
e.g. queue, stack, or just bag.

Conclusion and Links

- Fair synchronization primitives can be simple and fast
- Stronger guarantees \Rightarrow more complicated code

○

Conclusion and Links

- Fair synchronization primitives can be simple and fast
- Stronger guarantees \Rightarrow more complicated code

- Kotlin Coroutines project
github.com/Kotlin/kotlinx.coroutines
- The experiments (sqs-experiments branch)
github.com/Kotlin/kotlinx.coroutines/tree/segment-queue-synchronizer

Questions?